

AE 06: Joining country populations with continents

Jessica St Jean

Goal

Our ultimate goal in this application exercise is to create a bar plot of total populations of continents, where the input data are:

1. Countries and populations
2. Countries and continents

```
library(tidyverse) # for data wrangling and visualization  
library(scales)    # for pretty axis breaks
```

Data

Countries and populations

These data come from [The World Bank](#) and reflect population counts as of 2022.

```
population <- read_csv("data/world-pop-2022.csv")
```

Let's take a look at the data.

```
population
```

```
# A tibble: 217 x 3  
  country             year population  
  <chr>              <dbl>     <dbl>  
1 Afghanistan        2022     41129.  
2 Albania            2022     2778.
```

```

3 Algeria          2022    44903.
4 American Samoa  2022     44.3
5 Andorra          2022     79.8
6 Angola           2022    35589.
7 Antigua and Barbuda 2022    93.8
8 Argentina        2022    46235.
9 Armenia          2022    2780.
10 Aruba           2022     106.
# i 207 more rows

```

Continents

These data come from [Our World in Data](#).

```
continents <- read_csv("data/continents.csv")
```

Let's take a look at the data.

```
continents
```

```

# A tibble: 285 x 4
  entity            code   year continent
  <chr>           <chr> <dbl> <chr>
  1 Abkhazia        OWID_ABK 2015 Asia
  2 Afghanistan     AFG    2015 Asia
  3 Akrotiri and Dhekelia OWID_AKD 2015 Asia
  4 Aland Islands   ALA    2015 Europe
  5 Albania          ALB    2015 Europe
  6 Algeria          DZA    2015 Africa
  7 American Samoa  ASM    2015 Oceania
  8 Andorra          AND    2015 Europe
  9 Angola           AGO    2015 Africa
 10 Anguilla         AIA    2015 North America
# i 275 more rows

```

Exercises

- Think out loud:

- Which variable(s) will we use to join the population and continents data frames?

Comment: The country and entity variables.

- We want to create a new data frame that keeps all rows and columns from `population` and brings in the corresponding information from `continents`. Which join function should we use?

Comment: We would use `left_join` for this.

- **Demo:** Join the two data frames and name assign the joined data frame to a new data frame `population_continents`.

```
population_continents <- population |>
  left_join(continents, by = join_by(country == entity))
```

- **Demo:** Take a look at the newly created `population_continents` data frame. There are some countries that were not in `continents`. First, identify which countries these are (they will have NA values for `continent`).

```
population_continents |>
  filter(is.na(continent))
```

```
# A tibble: 6 x 6
  country                  year.x population code  year.y continent
  <chr>                   <dbl>      <dbl> <chr>  <dbl> <chr>
1 Congo, Dem. Rep.        2022      99010. <NA>     NA <NA>
2 Congo, Rep.             2022      5970.  <NA>     NA <NA>
3 Hong Kong SAR, China   2022      7346.  <NA>     NA <NA>
4 Korea, Dem. People's Rep. 2022      26069. <NA>     NA <NA>
5 Korea, Rep.              2022      51628. <NA>     NA <NA>
6 Kyrgyz Republic          2022      6975.  <NA>     NA <NA>
```

- **Demo:** All of these countries are actually in the `continents` data frame, but under different names. So, let's clean that data first by updating the country names in the `population` data frame in a way they will match the `continents` data frame, and then joining them, using a `case_when()` statement in `mutate()`. At the end, check that all countries now have continent information.

```
population_continents <- population |>
  mutate(country = case_when(
    country == "Congo, Dem. Rep." ~ "Democratic Republic of Congo",
    country == "Congo, Rep." ~ "Congo",
    country == "Hong Kong SAR, China" ~ "Hong Kong",
    country == "Korea, Dem. People's Rep." ~ "North Korea",
```

```

country == "Korea, Rep." ~ "South Korea",
country == "Kyrgyz Republic" ~ "Kyrgyzstan",
.default = country
)
) |>
left_join(continents, by = join_by(country == entity))

population_continent |>
  filter(is.na(continent))

```

```

# A tibble: 0 x 6
# i 6 variables: country <chr>, year.x <dbl>, population <dbl>, code <chr>,
#   year.y <dbl>, continent <chr>

```

- **Think out loud:** Which continent do you think has the highest population? Which do you think has the second highest? The lowest?

Comment: I think that the highest is Asia, second will probably be Africa and the lowest, if it included in this data set, would be Antarctica or if not it would likely be Oceania.

- **Demo:** Create a new data frame called `population_summary` that contains a row for each continent and a column for the total population for that continent, in descending order of population. Note that the function for calculating totals in R is `sum()`.

```

population_summary <- population_continent |>
  group_by(continent) |>
  summarize(total_pop = sum(population)) |>
  arrange (desc(total_pop))

population_summary

```

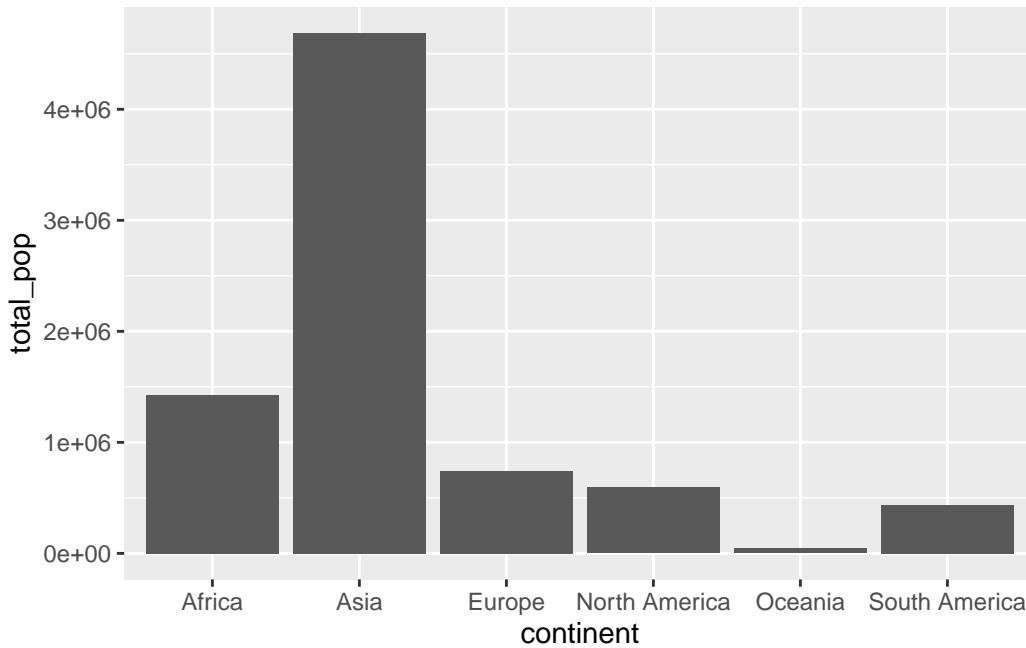
```

# A tibble: 6 x 2
  continent    total_pop
  <chr>        <dbl>
1 Asia          4685922.
2 Africa        1424831.
3 Europe         740607.
4 North America 594415.
5 South America 437233.
6 Oceania       44752.

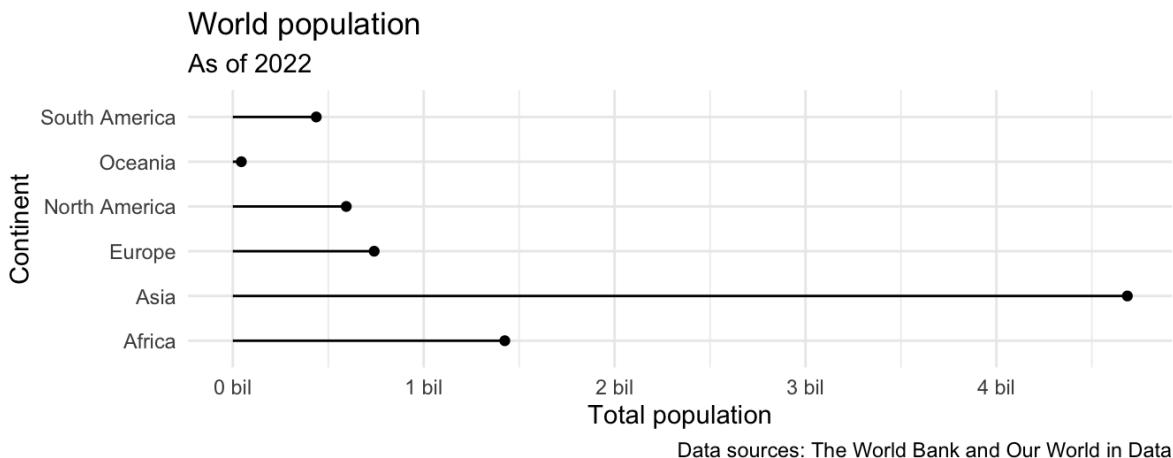
```

- **Your turn:** Make a bar plot with total population on the y-axis and continent on the x-axis, where the height of each bar represents the total population in that continent.

```
ggplot(population_summary, aes (x = continent, y= total_pop)) +
  geom_col()
```



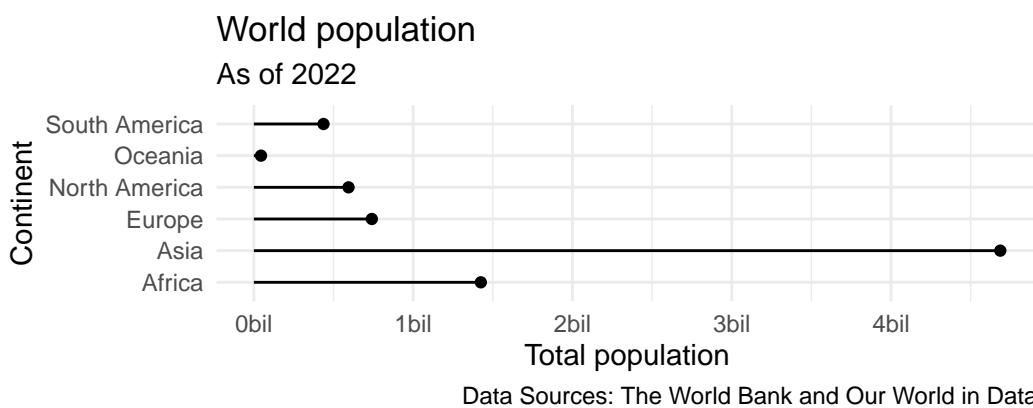
- **Your turn:** Recreate the following plot, which is commonly referred to as a **lollipop plot**. Hint: Start with the points, then try adding the segments, then add axis labels and caption, and finally, as a stretch goal, update the x scale (which will require a function we haven't introduced in lectures or labs yet!).



```

ggplot(population_summary, aes()) +
  geom_point(aes(x= total_pop, y = continent)) +
  geom_segment(aes(y= continent, x = 0, xend = total_pop)) +
  scale_x_continuous(labels = label_number(scale = 1/1000000, suffix = "bil")) +
  theme_minimal()+
  labs(
    x = "Total population",
    y = "Continent",
    title = "World population",
    subtitle = "As of 2022",
    caption = "Data Sources: The World Bank and Our World in Data"
)

```



- **Think out loud:** What additional improvements would you like to make to this plot.

Comment: Orient the graph so that continents are on the x-axis, add some clearer labels, re-order the continents in order of highest to lowest population.